

1.Code Optimization – “Beat the Compiler”

(a)

(i) gcc -O2 -o loop_performance loop_performance.c

- ./loop_performance 10000000
run 10 times, the shortest time is 23.576000 milliseconds
- ./loop_performance 100000000
run 10 times, the shortest time is 232.444000 milliseconds

(ii) gcc -O3 -o loop_performance loop_performance.c

- ./loop_performance 10000000
run 10 times, the shortest time is 19.212000 milliseconds
- ./loop_performance 100000000
run 10 times, the shortest time is 187.759000 milliseconds

(b)

- I was running the code on a VM provided by duke vcm and the processor arch is x86_64, cpu freq is 2.70GHz, OS is Ubuntu18.

(c)

- Original

```
void do_loops(int *a, int *b, int *c, int N)
{
    int i;
    for (i=N-1; i>=1; i--) {
        a[i] = a[i] + 1;
    }
    for (i=1; i<N; i++) {
        b[i] = a[i+1] + 3;
    }
    for (i=1; i<N; i++) {
        c[i] = b[i-1] + 2;
    }
}
```

- Loop Fusion

By doing so, the optimized -O3 version is slower than original -O3, and the optimized -O2 version is about the same with the original -O2. Therefore, I think the parallellism is kind of hampered in the optimized version.

- Loop unrolling

By doing so, the number of loops is fewer by seeing the assembly code.

- Loop reversal

I reversed the loop order of the loops containing b and c. This did not work well. The -O2 version was about the same while the -O3 version is even slower.

- Final code (Combined method)

```
void do_loops(int *a, int *b, int *c, int N)
{
    int i;
    b[N-1] = a[N] + 3;
    for (i=N-1; i>=2; i--) {
        a[i] = a[i] + 1;
        b[i-1] = a[i] + 3;
        c[i] = b[i-1] + 2;
    }
    c[1] = b[0] + 2;
    a[1] = a[1] + 1;
}
```


I expected this to work well but it didn't. This could beat the -O2 compiler but not the -O3.

(d)

Generally, I've tried a lot of methods and their combinations, but I cannot beat -O3 compiler. The -O2 compiler is easier to beat. Later I found that this may have something to do with the OS version and gcc version. I'm using Ubuntu18 with gcc 7.5.0, but the same code compiled on a VM with Ubuntu20.04 and gcc 9.1.3 beats the -O3 compiler successfully.

2.Dependence Analysis

(a)



(b)

- Loop independent dependencies:

$S1[i,j] \rightarrow T S3[i,j]$

$S1[i,j] \rightarrow A S2[i,j]$

- Loop-carried dependencies:

$S1[i,j] \rightarrow T S2[i+1,j+1]$

$S1[i,j] \rightarrow A \ S1[i+1,j-1]$


$S3[i-1,j] \rightarrow T \ S1[i,j]$

$S4[i-1,j+1] \rightarrow T \ S4[i,j]$

$S3[i-2,j] \rightarrow T \ S2[i,j]$

(c)

- Loop-carried Dependence Grap



3.Function In-lining and Performance

(a)

(i) with attribute noline

```
int add (int a, int b) __attribute__ ((__noinline__));  
  
int add (int a, int b) {  
    return (a+b);  
}
```

- run 10 times, the shortest time is 254.035 milliseconds

(ii) with attribute `always_inline`

```
int add (int a, int b) __attribute__ ((__always_inline__));

int add (int a, int b) {
    return (a+b);
}
```

- run 10 times, the shortest time is 104.934 milliseconds

(b)

noinline assembly code

- loop

```

e12:  e8 19 fd ff ff      callq  b30 <gettimeofday@plt>
e17:  44 89 e0             mov     %r12d,%eax
e1a:  4c 8b 54 24 50      mov     0x50(%rsp),%r10
e1f:  4c 8b 4c 24 30      mov     0x30(%rsp),%r9
e24:  4c 8b 44 24 70      mov     0x70(%rsp),%r8
e29:  48 8d 0c 85 04 00 00 lea     0x4(,%rax,4),%rcx
e30:  00
e31:  31 d2              xor     %edx,%edx
e33:  0f 1f 44 00 00      nopl    0x0(%rax,%rax,1)
e38:  41 8b 34 12         mov     (%r10,%rdx,1),%esi
e3c:  41 8b 3c 11         mov     (%r9,%rdx,1),%edi
e40:  e8 eb 04 00 00      callq   1330 <_Z3addii>
e45:  41 89 04 10         mov     %eax,(%r8,%rdx,1)
e49:  48 83 c2 04         add     $0x4,%rdx
e4d:  48 39 ca           cmp     %rcx,%rdx
e50:  75 e6             jne     e38 <main+0x2c8>
e52:  31 f6             xor     %esi,%esi
e54:  48 89 ef         mov     %rbp,%rdi
e57:  e8 d4 fc ff ff      callq   b30 <gettimeofday@plt>
```

- `add()` function

```

00000000000001330 <_Z3addii>:
1330:  8d 04 37           lea     (%rdi,%rsi,1),%eax
1333:  c3                retq
1334:  66 2e 0f 1f 84 00 00 nopw    %cs:0x0(%rax,%rax,1)
133b:  00 00 00
133e:  66 90             xchg    %ax,%ax
```

- It is clear that for noinline code, every time the code calls `add()` function, in assembly code it actually branch to somewhere else for the code.

always inline assembly code

```

e12:  e8 19 fd ff ff      callq  b30 <gettimeofday@plt>
e17:  48 8b 7c 24 50      mov     0x50(%rsp),%rdi
e1c:  48 8b 74 24 70      mov     0x70(%rsp),%rsi
e21:  4c 8b 44 24 30      mov     0x30(%rsp),%r8
e26:  48 8d 47 10         lea     0x10(%rdi),%rax
e2a:  48 8d 4e 10         lea     0x10(%rsi),%rcx
e2e:  48 39 c6           cmp     %rax,%rsi
e31:  0f 93 c2         setae   %dl
e34:  48 39 cf         cmp     %rcx,%rdi
e37:  0f 93 c0         setae   %al
e3a:  09 c2             or      %eax,%edx
e3c:  49 8d 40 10         lea     0x10(%r8),%rax
e40:  48 39 c6           cmp     %rax,%rsi
e43:  0f 93 c0         setae   %al
e46:  49 39 c8         cmp     %rcx,%r8
e49:  0f 93 c1         setae   %cl
e4c:  09 c8             or      %ecx,%eax
e4e:  84 c2             test    %al,%dl
e50:  0f 84 7e 03 00 00   je      11d4 <main+0x664>
e56:  83 fb 08         cmp     $0x8,%ebx
e59:  0f 86 75 03 00 00   jbe     11d4 <main+0x664>
e5f:  48 89 f9         mov     %rdi,%rcx
```

e62:	48 c1 e9 02	shr	\$0x2,%rcx
e66:	48 f7 d9	neg	%rcx
e69:	83 e1 03	and	\$0x3,%ecx
e6c:	0f 84 5a 03 00 00	je	11cc <main+0x65c>
e72:	41 8b 00	mov	(%r8),%eax
e75:	03 07	add	(%rdi),%eax
e77:	83 f9 01	cmp	\$0x1,%ecx
e7a:	89 06	mov	%eax,(%rsi)
e7c:	0f 84 b0 03 00 00	je	1232 <main+0x6c2>
e82:	41 8b 40 04	mov	0x4(%r8),%eax
e86:	03 47 04	add	0x4(%rdi),%eax
e89:	83 f9 02	cmp	\$0x2,%ecx
e8c:	89 46 04	mov	%eax,0x4(%rsi)
e8f:	0f 84 dc 03 00 00	je	1271 <main+0x701>
e95:	41 8b 40 08	mov	0x8(%r8),%eax
e99:	03 47 08	add	0x8(%rdi),%eax
e9c:	41 b9 03 00 00 00	mov	\$0x3,%r9d
ea2:	89 46 08	mov	%eax,0x8(%rsi)
ea5:	41 89 db	mov	%ebx,%r11d
ea8:	31 c0	xor	%eax,%eax
ea:	31 d2	xor	%edx,%edx
eac:	41 29 cb	sub	%ecx,%r11d
eaf:	89 c9	mov	%ecx,%ecx
eb1:	48 c1 e1 02	shl	\$0x2,%rcx
eb5:	45 89 da	mov	%r11d,%r10d
eb8:	4c 8d 34 0f	lea	(%rdi,%rcx,1),%r14
ebc:	4d 8d 2c 08	lea	(%r8,%rcx,1),%r13
ec0:	41 c1 ea 02	shr	\$0x2,%r10d
ec4:	48 01 f1	add	%rsi,%rcx
ec7:	66 0f 1f 84 00 00 00	nopw	0x0(%rax,%rax,1)
ece:	00 00		
ed0:	f3 41 0f 6f 44 05 00	movdqu	0x0(%r13,%rax,1),%xmm0
ed7:	83 c2 01	add	\$0x1,%edx
eda:	66 41 0f fe 04 06	padd	(%r14,%rax,1),%xmm0
ee0:	0f 11 04 01	movups	%xmm0,(%rcx,%rax,1)
ee4:	48 83 c0 10	add	\$0x10,%rax
ee8:	41 39 d2	cmp	%edx,%r10d
eeb:	77 e3	ja	ed0 <main+0x360>
eed:	44 89 da	mov	%r11d,%edx
ef0:	83 e2 fc	and	\$0xffffffff,%edx
ef3:	41 39 d3	cmp	%edx,%r11d
ef6:	42 8d 04 0a	lea	(%rdx,%r9,1),%eax
efa:	74 70	je	f6c <main+0x3fc>
efc:	48 63 d0	movslq	%eax,%rdx
eff:	41 8b 0c 90	mov	(%r8,%rdx,4),%ecx
f03:	03 0c 97	add	(%rdi,%rdx,4),%ecx
f06:	89 0c 96	mov	%ecx,(%rsi,%rdx,4)
f09:	8d 50 01	lea	0x1(%rax),%edx
f0c:	39 d3	cmp	%edx,%ebx
f0e:	7e 5c	jle	f6c <main+0x3fc>
f10:	48 63 d2	movslq	%edx,%rdx
f13:	41 8b 0c 90	mov	(%r8,%rdx,4),%ecx
f17:	03 0c 97	add	(%rdi,%rdx,4),%ecx
f1a:	89 0c 96	mov	%ecx,(%rsi,%rdx,4)
f1d:	8d 50 02	lea	0x2(%rax),%edx
f20:	39 d3	cmp	%edx,%ebx
f22:	7e 48	jle	f6c <main+0x3fc>
f24:	48 63 d2	movslq	%edx,%rdx
f27:	41 8b 0c 90	mov	(%r8,%rdx,4),%ecx
f2b:	03 0c 97	add	(%rdi,%rdx,4),%ecx
f2e:	89 0c 96	mov	%ecx,(%rsi,%rdx,4)
f31:	8d 50 03	lea	0x3(%rax),%edx
f34:	39 d3	cmp	%edx,%ebx
f36:	7e 34	jle	f6c <main+0x3fc>
f38:	48 63 d2	movslq	%edx,%rdx
f3b:	41 8b 0c 90	mov	(%r8,%rdx,4),%ecx
f3f:	03 0c 97	add	(%rdi,%rdx,4),%ecx
f42:	89 0c 96	mov	%ecx,(%rsi,%rdx,4)
f45:	8d 50 04	lea	0x4(%rax),%edx
f48:	39 d3	cmp	%edx,%ebx
f4a:	7e 20	jle	f6c <main+0x3fc>
f4c:	48 63 d2	movslq	%edx,%rdx
f4f:	83 c0 05	add	\$0x5,%eax
f52:	41 8b 0c 90	mov	(%r8,%rdx,4),%ecx
f56:	03 0c 97	add	(%rdi,%rdx,4),%ecx
f59:	39 c3	cmp	%eax,%ebx
f5b:	89 0c 96	mov	%ecx,(%rsi,%rdx,4)
f5e:	7e 0c	jle	f6c <main+0x3fc>
f60:	48 98	cltq	


```
f62:    41 8b 14 80          mov    (%r8,%rax,4),%edx
f66:    03 14 87             add    (%rdi,%rax,4),%edx
f69:    89 14 86             mov    %edx,(%rsi,%rax,4)
f6c:    31 f6               xor    %esi,%esi
f6e:    48 89 ef             mov    %rbp,%rdi
f71:    e8 ba fb ff ff       callq  b30 <gettimeofday@plt>
```

- While in `always_inline` code, the loop was flattened and execute many times, without branching soemwhere else, which significantly reduced the instruction number executed.

(c)

- Yes, it matched my expectation. The time cost of code compiled inline should be less than the one compiled noinline, since the total instructions executed was less.

(d)

- For no attribute on `add()` function, out of 10 times, the shortest time is 104.594 milliseconds. I would give a guess that the compiler was in-lining the `add()` function by default.

4. Loop transformations I

Original Code

```
int a[N][4];
int rand_number = rand();
for (i=0; i<4; i++) {
    threshold = 2.0 * rand_number;
    for (j=0; j<N; j++) {
        if (threshold < 4) {
            sum = sum + a[j][i];
        } else {
            sum = sum + a[j][i] + 1;
        }
    }
}
```

Loop Invariant Hoisting

```
int a[N][4];
int rand_number = rand();
threshold = 2.0 * rand_number;
for (i=0; i<4; i++) {
    for (j=0; j<N; j++) {
        if (threshold < 4) {
            sum = sum + a[j][i];
        } else {
            sum = sum + a[j][i] + 1;
        }
    }
}
```

Loop unswitching

```
int a[N][4];
int rand_number = rand();
threshold = 2.0 * rand_number;
if (threshold < 4) {
    for (i=0; i<4; i++) {
        for (j=0; j<N; j++) {
            sum = sum + a[j][i];
        }
    }
} else {
    for (i=0; i<4; i++) {
        for (j=0; j<N; j++) {
            sum = sum + a[j][i] + 1;
        }
    }
}
```

Loop Interchange

```
int a[N][4];
int rand_number = rand();
threshold = 2.0 * rand_number;
if (threshold < 4) {
    for (j=0; j<N; j++) {
        for (i=0; i<4; i++) {
            sum = sum + a[j][i];
        }
    }
} else {
    for (j=0; j<N; j++) {
        for (i=0; i<4; i++) {
            sum = sum + a[j][i] + 1;
        }
    }
}
```

Loop Unrolling

```
int a[N][4];
int rand_number = rand();
threshold = 2.0 * rand_number;
if (threshold < 4) {
    for (j=0; j<N; j++) {
        sum = sum + a[j][0];
        sum = sum + a[j][1];
        sum = sum + a[j][2];
        sum = sum + a[j][3];
    }
} else {
    for (j=0; j<N; j++) {
        sum = sum + a[j][0];
        sum = sum + a[j][1];
        sum = sum + a[j][2];
        sum = sum + a[j][3];
    }
}
```

5. Loop transformations II

(a)

- Not safe. Originally, there is a loop-carried output dependency from S1[i] to S3[i-1] and a loop-carried anti dependency from S2[i] to S3[i-1], while after the transformation, there is a loop-carried output dependency from S3[i] to S1[i+1] and a loop-carried true dependency from S3[i] to S2[i+1].

(b)

- Not safe. Originally, the outermost loop carries a loop-carried anti dependency from S1[i][j] to S1[i+1][j-1]. Here i < i' and j > j', so it's not safe.

(c)

- Safe.